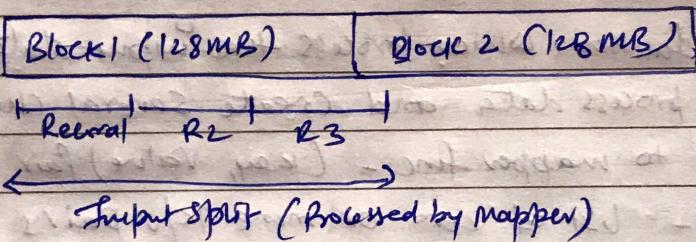


## Input Split

Date: \_\_\_\_\_ Page: 11

- \* logical representation of data
- \* unit of work that contains a single map task.
- \* represent data which is processed by an individual mapper.
- \* Split  $\xrightarrow{\text{divide}}$  records  $\xrightarrow{\text{Process}}$  Mapper  
(key, value)
- \* Doesn't contain the data, reference to data
- \* InputFormat create InputSplit and divide into record
  - ↳ file input format - break into 128 MB chunks
  - ↳ mapred.min.Split.Size (mapred.xml) (default)



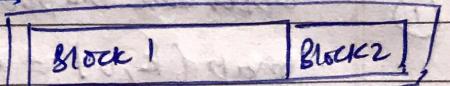
- \* Using starting record in block + byte offset - get complete record even if it spans the block boundaries

## Input Split

- \* logical rep of data
- \* appr equal to block size by default - can configure
- \* file split into 2 blocks
- \* No of map tasks (mapper) are equal to no of extramarks

## HDFS Block

- \* Physical rep of data
- \* default size - 128 MB



(logical grouping of blocks)

Input Split

## Record Reader

- \* Load data from source and convert into key-value pair suitable for reading by mapper.
- \* uses start / end byte position - generating K/V pair.
- \* Communicate with Input Split until file reading complete works
- \* Iterate over records
- \* map task uses one record to generate key-value pair - passed to map function

## Mapper

- \* A function which processes the input data
  - \* It processes data and create several chunks of data
  - \* I/O to mapper fine - (key, value) pair
  - \* O/P - list of zero or more K/V pairs
  - \* May use or completely ignore the input keys
- eg - Standard pattern - read one line at a time
- key byte offset , value - Content of line
  - Key is considered irrelevant
  - If mapper write anything out , the output must be in the form of key/value pairs

(1)

### Upper Case Mapper

$\text{map}(K, V) = \text{emit}(K.\text{toUpper}, V.\text{toUpper})$

(2)

### Explode

$\text{map}(K, V) = \text{foreach char } c \text{ in } V \text{ emit }(K, c)$

extra marks

- (3) filter  
 $\text{map}(k, v) = \text{if } (\text{isprime}(v)) \text{ then emit}(k, v)$
- (4) changing keyspace  
 $\text{map}(k, v) = \text{emit}(\text{v.length}(), v)$
- (5) Identity mapper  
 $\text{map}(k, v) = \text{emit}(k, v)$

## Reducer

- \* Process o/p of the mapper
- \* After processing produces a new set of o/p.
- \* Take intermediate key-value pair as I/P and runs reducer function on each of them.  
 (e.g. aggregate, ~~filter~~, combine the data)
- \* Reducers run in parallel - as independent
- \* User can decide the no. of reducer (def - 1)
- \* HDFS stores the o/p data
- \* Job.setNumReduceTasks(1st) - Set no. of reducer

### 3 Phases

- 1) Shuffle - Reducer copies the sorted out put from each mapper using HTTP across the n/w
  - 2) Sort - Framework merge sort reduce input by keys (diff mapper - may have same key)
  - 3) Secondary Sort
  - 3) Reduce - reduce method is called on each key in the sorted input  
 the o/p of reduce task to RecordWriter
- extramarks** ➤ o/p is not sorted

## eg 1) Sum reducer

$\text{reduce}(K, \text{vals}) = \text{Sum} = 0;$

foreach int i in vals:

$\text{Sum} += i$

emit  $(K, \text{Sum})$

## 2) Average reducer

$\text{reduce}(K, \text{val}) = \text{Sum} = 0, \text{Counter} = 0,$

foreach int i in Val:

$\text{Sum} += i;$  Counter += 1;

emit  $(K, \text{Sum}/\text{Counter})$

## 3) Identity reducer

$\text{reduce}(K, \text{val}) = \text{for each } v \text{ in val :}$

emit  $(K, v)$

Speculative Execution

- \* MR breaks job - task - run in parallel to reduce overall exec time
- \* Sensitive to Slow task - slowdown type of job
- \* Reason - H/W degradation - soft mis configuration
- \* Hadoop doesn't try to fix, it detects and runs
- \* It launches another parallel / backup task for each task that is performing slow than expected on faster nodes
- \* these are called Speculative tasks.

## Assignment - Benefit / Limitation (Ques of S.E)

Date: \_\_\_\_\_ Page: 15

- \* Task Complete - inform job tracker.
- \* Whenever copy finishes first - definitive copy
- \* Kill other task (Slow)
- \* Enabled by default
- \* - mapred.map.task.speculative.execution [Config]
- \* mapred.reduce.task.speculative.execution