

Resource Request

- Request for containers - amount of compute res. and locality constraint
- Locality is critical - to use cluster bandwidths efficiently
- If locality cannot be met - YARN starts container on node in same rack or any node in cluster
- YARN application can make res. req.
 - ↳ all request up front (Spark)
 - ↳ Dynamic approach (MR - Map → Reduce)

Application lifespan

- Can be short lived app (few sec) or long running app (runs for days or months)
- Categorize app - how they map to the job user run.
 - ↳ one application per user job (MR)
 - ↳ one app per workflow or user session (unrelated jobs)
 - ↳ more efficient | container reused b/w jobs | Cache intermediate data b/w jobs (Spark)
 - ↳ long running app shared by diff users
 - ↳ coordination role
 - ↳ Apache Glider - has AM for launching other app on cluster. (Impala)

YARN Applications Interested in

- Running (DAG) Job - Spark / Tez
- Stream processing - Spark Streaming, Samza / Storm
- Simplify process of building YARN - Apache Girder -
Possible to run existing dist app on YARN - Different user can run diff version of same app (HBase)

MapReduce 1 | YARN -

1 - JobTracker	- RM, AM, Timeline Server
2 - TaskTracker	- Node Manager
3 - Slot	- Container

Timeline Server - Store application history

4 - Batch Processing	- Real time, Interactive
5 - 1 nameSpace	- Multiple namespace
6 - Only MR app	- MapReduce and non MR app
7 - Average C/Ses optimization (Fixed MRSlot)	- Excellent cluster res. optimization (Central)

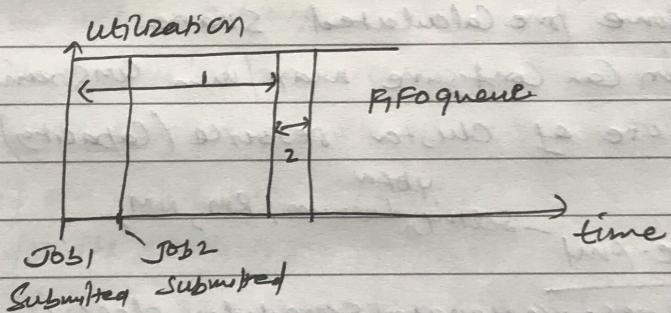
YARN Scheduling

Date: _____ Page: 27

- Scheduler is a component of RM
- Responsible for allocating resources to app.
- Scheduler
 - ↳ perform no monitoring or tracking status of app
 - ↳ offer no guarantees about restarting failed task

1) FIFO Scheduler

- First in first out
- Jobs run in the order of submission
- Job submitted first will get priority to execute
- queue based scheduler
- Simple, does not guarantee performance effectively



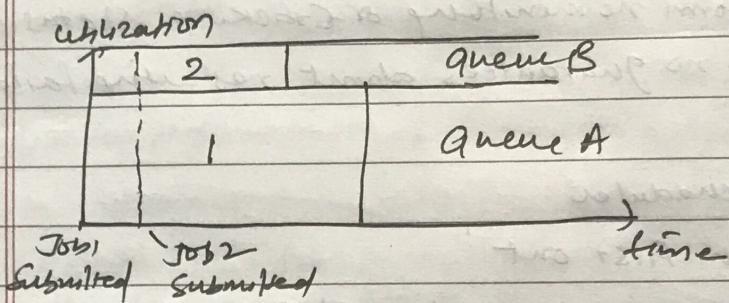
- Simple to understand - not needing any configuration
- not suitable for shared cluster.
- large app will use all cluster resources - small one to wait

2) Capacity Scheduler

- Allow app to share resource
- Job queues
- separate queue - allow small job to start as soon as it is submitted

extramarks

- As Capacity queue is reserved for jobs - at the cost of overall cluster utilization
- large job finishes later than FIFO



- allow user to share resources - User or group of user assigned a certain capacity of cluster.
- Tenable - admin config one or more queue with some pre calculated share
- Admin can configure max/min constraint on the use of cluster resource (Capacity) on each queue

`YARN-site.xml` - setting, domain: RM, NM,

name - `Yarn.resourcemanager.scheduler.class`

value - `org.apache.hadoop.yarn.server.resourcemanager.Scheduler.CapacityScheduler` (Capacity)

/ `FifoScheduler` (FIFO)

/ `FairScheduler` (Fair)

Capacity-Scheduler.xml

name - `Yarn.scheduler.Capacity.root-queues`

value - `prod, dev`

name - yarn.scheduler.capacity.root.dev.queue
 value - Job1, Job2
 name - yarn.scheduler.capacity.root.prod.capacity
 value - 40
 name - yarn.scheduler.capacity.root.dev.Capacity
 value - 60
 name - yarn.scheduler.capacity.root.dev.maximumCapacity
 value - 75

Fair Scheduling

- Enable memory intensive app to share cluster resource in efficient way.
- Allocation of resources such that all applications get, on average an equal share of resource
- if one app is running - might request all resource for its execution if needed.
- if other app are submitted policy distribute free resource among app such that each app get a fairly equal share of resource.
- allow preemption - RM request container back from AM depends of Job Configuration
- Adv - Every queue get min share of cluster res.
- App waiting in queue - get min resource share
- Resources are more than for app - excess amount would be distributed.

Utilization